

Quadratic Optimization

This handout covers some results on constrained nonlinear optimization. Suppose x is a variables vector of n dimensions: $x = (x_1, \dots, x_n)$. We consider simple polynomial function of x that involve at most second-order terms. Matrix and vector notation considerably simplify the expressions. Here is a reminder of several formulas, in terms of the vectors x, y and the matrices A, B :

The inner product of x and y	$x'y = x_1y_1 + \dots + x_ny_n$
The squared norm of x	$ x ^2 = x'x = x_1^2 + \dots + x_n^2$
The most general quadratic form	$x'Ax + b'x + c$
Remainder	$x'Ax = x'(Ax) = \sum_{ij} a_{ij}x_ix_j$
Positive semidefinite matrix B	$x'Bx \geq 0$ whenever $x \neq 0$
Gradients (derivatives)	$f(x) = x'Bx + b'x + c$, B pos semidef, then $\nabla f(x) = 2Bx + b$

Unconstrained optimization

Here we would like to compute the minimum of $f(x)$. It can be shown that a quadratic $f(x) = x'Bx + b'x + c$ has a minimum if and only if B is positive semidefinite. It has a unique minimum if and only if B is positive definite. In both cases the minimum is values of x that zero out the gradient. That is, the solutions of:

$$Bx = -\frac{b}{2}$$

Example 1

Find the minimum of

$$f(x_1, x_2) = x_1^2 + 2x_2^2 + 2x_1x_2 + 2x_1 - 6x_2 + 4$$

With $x = (x_1, x_2)$ we have:

$$f(x) = x' \begin{pmatrix} 1 & 1 \\ 1 & 2 \end{pmatrix} x + (2, -6)x + 4$$

Therefore, if there is a minimum it is the solution of:

$$\begin{pmatrix} 1 & 1 \\ 1 & 2 \end{pmatrix} x = -(2, -6)/2 = (-1, 3)$$

The solution is $(-5, 4)$ so that if a minimum exists it is at $x_1 = -5, x_2 = 4$. To argue that it is indeed a minimum we still need to show that the matrix is positive definite. One way of showing that a matrix B is positive semidefinite is to show that there is a matrix G such that:

$$B = G'G$$

In our case with $B = \begin{pmatrix} 1 & 1 \\ 1 & 2 \end{pmatrix}$ we can take $G = \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix}$.

Equality linear constraints

Here we would like to compute the minimum of $f(x)$ where x is also subject to linear equality constraints. A linear equality constraint can be written as $c'x = d$ where c is a vector and d is a scalar. When there are k constraints we write them as $c'_i x = d_i$.

There are several methods to reduce this to the case of unconstrained optimization. The most basic one is the method of Lagrange multipliers. In order to minimize $f(x)$ under the above equality constraints we define a Lagrangian multiplier α_i for the i th constraint and form the Lagrangian function $L(x, \alpha_1, \dots, \alpha_k)$ as follows:

$$L(x, \alpha_1, \dots, \alpha_k) = f(x) + \sum_{i=1}^k \alpha_i (c'_i x - d_i)$$

We can now treat L as a function of $n + k$ variables and use the method of unconstrained minimization to find the unconstrained minimum of L . (Observe that L is quadratic.)

Example 2

Minimize $f(x_1, x_2, x_3) = x_1^2 + x_2^2 + x_3^2$ under the two constraints: $x_2 - x_1 = 1$, $x_3 = 1$.

This is an easy case to solve without the Lagrange technique. It is easy to see by direct substitution that $f = x_1^2 + (x_1 + 1)^2 + 1$ and by taking derivatives and equating to 0 the minimum is at $x_1 = -1/2$. Using the Lagrange technique:

$$L(x_1, x_2, x_3, \alpha_1, \alpha_2) = x_1^2 + x_2^2 + x_3^2 + \alpha_1(x_2 - x_1 - 1) + \alpha_2(x_3 - 1)$$

The derivative with respect to x_1 gives:	$2x_1 - \alpha_1 = 0$
The derivative with respect to x_2 gives:	$2x_2 + \alpha_1 = 0$
The derivative with respect to x_3 gives:	$2x_3 + \alpha_2 = 0$
The derivative with respect to α_1 gives:	$x_2 - x_1 - 1 = 0$
The derivative with respect to α_2 gives:	$x_3 - 1 = 0$

Treating this as a system of 5 equations with 5 unknowns we get the following solution:

$$x_1 = -1/2, \quad x_2 = 1/2, \quad x_3 = 1, \quad \alpha_1 = -1, \quad \alpha_2 = -2$$

The dual problem

There is another way of using the Lagrangian to solve the constrained optimization problem. The following result can be proved:

$$\min_{x \text{ subject to the equality constraints}} f(x) = \max_{\alpha_1, \dots, \alpha_k} \min_x L(x, \alpha_1, \dots, \alpha_k)$$

To use this, we begin by treating the Lagrangian multipliers as constants, minimizing the Lagrangian with respect to x . This is done by taking the derivatives only with respect to x . The resulting system of linear equations can be used to solve for the variables x in terms of the Lagrangian multipliers. This is then substituted back into the Lagrangian, transforming it into a function of only the Lagrangian multipliers. This function is called **the dual problem**, and it can then be solved using unconstrained techniques.

Example 3

To convert the problem of Example 2 into a dual form we only need to take the derivatives of the Lagrangian with respect to x_1, x_2, x_3 As above:

The derivative with respect to x_1 gives:	$2x_1 - \alpha_1 = 0$
The derivative with respect to x_2 gives:	$2x_2 + \alpha_1 = 0$
The derivative with respect to x_3 gives:	$2x_3 + \alpha_2 = 0$

Solving for x_1, x_2, x_3 in terms of α_1, α_2 we have:

$$x_1 = \alpha_1/2, \quad x_2 = -\alpha_1/2, \quad x_3 = -\alpha_2/2 \quad (1)$$

Substituting back into the Lagrangian gives:

$$L(\alpha_1, \alpha_2) = -\frac{1}{2}\alpha_1^2 - \frac{1}{4}\alpha_2^2 - \alpha_1 - \alpha_2$$

The dual problem requires that we maximize this expression. The maximizing α_1, α_2 are: $\alpha_1 = -1, \alpha_2 = -2$, and the minimizing x_1, x_2, x_3 of the original (primal) problem are computed from (1).

Example 4

Consider the n dimensional hyperplane (generalization of a line in 2D) $w'x = s$. In 2D it is the line $w_1x_1 + w_2x_2 = s$. Compute the distance of this hyperplane from the origin.

A point x on the hyperplane has distance of $|x|$ from the origin. We are looking for the point x that minimizes $|x|^2$ subject to the constraint $w'x = s$. The Lagrangian:

$$L(x, \alpha) = |x|^2 + \alpha(w'x - s)$$

Computing the derivative with respect to x and equating to 0 we have:

$$2x + \alpha w = 0 \quad \Rightarrow \quad x = -\frac{\alpha}{2}w$$

Substituting into the Lagrangian we get the dual problem:

$$L(\alpha) = \frac{\alpha^2}{4}|w|^2 + \alpha(-\frac{\alpha}{2}w'w - s) = -\frac{\alpha^2}{4}|w|^2 - \alpha s$$

In terms of α this is a parabola with maximum at $\alpha = -2s/|w|^2$. This gives: $x = sw/|w|^2$ so that the distance of the hyperplane from the origin is $|s|/|w|$.

Linear inequality constraints

Here we would like to compute the minimum of $f(x)$ where x is subject to linear inequality constraints. A linear inequality constraint can be written as $c'x \leq d$ where c is a vector and d is a scalar. When there are k constraints we write them as $c'_i x \leq d_i$.

The Lagrangian in this case is the same as in the previous case:

$$L(x, \alpha_1, \dots, \alpha_k) = f(x) + \sum_{i=1}^k \alpha_i (c'_i x - d_i)$$

And the following dual theorem holds:

$$\min_{x \text{ subject to the inequality constraints}} f(x) = \max_{\alpha_1, \dots, \alpha_k \text{ subject to } \alpha_j \geq 0} \min_x L(x, \alpha_1, \dots, \alpha_k)$$

Formally, the problem

$$\min_{x \text{ subject to the inequality constraints}} f(x)$$

is called **the primal problem** and

$$\max_{\alpha_1, \dots, \alpha_k \text{ subject to } \alpha_j \geq 0} \min_x L(x, \alpha_1, \dots, \alpha_k)$$

is called **the dual problem**. Observe that here the dual problem is no longer trivial.

The following important result can be proved for the solution to the dual problem:

The Karush-Kuhn-Tucker complementarity condition:

$$\alpha_i (c'_i x - d_i) = 0 \quad \text{for } i = 1, \dots, k$$

This means that either the constraint holds as an **equality** constraint, or the corresponding α_i is 0.

Example 5

Minimize x^2 subject to $x \geq 1$ and $x \geq 2$. The minimizing value of x is clearly $x = 2$, with the first constraint inactive. We show how this can be deduced from the math describe above. First write the constraints as required:

$$\begin{aligned} -x + 1 &\leq 0 \\ -x + 2 &\leq 0 \end{aligned}$$

The Lagrangian:

$$L(x, \alpha_1, \alpha_2) = x^2 + \alpha_1(-x + 1) + \alpha_2(-x + 2)$$

Now compute the dual problem. Taking the derivative w.r.t. x and equating to 0:

$$2x - \alpha_1 - \alpha_2 = 0 \quad \Rightarrow \quad x = \frac{\alpha_1 + \alpha_2}{2}$$

Substituting back into the Lagrangian:

$$L(\alpha_1, \alpha_2) = \frac{(\alpha_1 + \alpha_2)^2}{4} + \alpha_1(1 - \frac{\alpha_1 + \alpha_2}{2}) + \alpha_2(2 - \frac{\alpha_1 + \alpha_2}{2}) = \alpha_1 + 2\alpha_2 - \frac{(\alpha_1 + \alpha_2)^2}{4}$$

The dual problem:

$$\max_{\alpha_1 \geq 0, \alpha_2 \geq 0} \alpha_1 + 2\alpha_2 - \frac{(\alpha_1 + \alpha_2)^2}{4}$$

The solution is $\alpha_1 = 0, \alpha_2 = 4$.

To see why, assume that the optimum is obtained for $\alpha_1 = a_1 > 0, \alpha_2 = a_2 \geq 0$. Then taking $\alpha_1 = 0, \alpha_2 = a_1 + a_2$ increases the dual, contradiction. Now if $\alpha_1 = 0$ the global max is at $\alpha_2 = 4$.

The dual problem is assumed to be solved by a “black box” software.